

NVIDIA Tesla P4 GPU - Test Report

Date: February 20, 2026 Prepared for: Carrie (Vendor) Prepared by: Alphonzo Jones Jr.

Device Information

| Field | Value |
|----------------|---|
| GPU Model | NVIDIA Tesla P4 |
| VRAM | 8 GB GDDR5 (7680 MiB) |
| Host System | HP Z240 Workstation |
| PCIe Bus ID | 00000000:02:00.0 |
| Driver Version | 570.211.01 |
| CUDA Version | 12.8 |
| OS | Ubuntu 24.04.3 LTS (Kernel 6.8.0-100-generic) |

Summary

The NVIDIA Tesla P4 GPU fails under compute load with **Xid 79 errors** ("GPU has fallen off the bus") and **PCIe RxErr** (receive errors on the PCIe bus). The GPU responds to basic queries at idle but cannot sustain any workload. The card was confirmed defective via gpu-burn stress testing on February 19, 2026. All GPU workloads have been migrated to CPU-only operation.

Test Results

Test 1: nvidia-smi at Idle - PASS

The GPU responds to nvidia-smi when no compute workload is running.

```
+-----+-----+-----+
| NVIDIA-SMI 570.211.01      Driver Version: 570.211.01      CUDA Version: 12.8 |
|-----+-----+-----+
| GPU  Name                  Persistence-M | Bus-Id      Disp.A | Volatile Uncorr. ECC |
| Fan  Temp     Perf          Pwr:Usage/Cap | Memory-Usage | GPU-Util  Compute M. |
|-----+-----+-----+-----+-----+-----+-----+-----+
| 0  Tesla P4                 Off  | 00000000:02:00.0 Off |          0 |
| N/A  57C     P8              8W / 75W |    0MiB /  7680MiB |     0%     Default |
+-----+-----+-----+-----+-----+-----+-----+-----+
```

Result: GPU is detected and reports healthy status at idle. Temperature 57C, power draw 8W (idle).

Test 2: LLM Inference Workload (Ollama) - FAIL

Deployed Ollama (LLM inference server) with `nvidia.com/gpu: 1` resource limit, loading a 3B-parameter model (llama3.2).

Observed behavior: - Model loads into GPU memory initially (~4.9 GB VRAM consumed) - After variable runtime under inference load, the GPU becomes unresponsive - Kernel logs report **NVIDIA Xid 79 error** - Pod events show: failed to get device handle from UUID: Unknown Error - Requires full node reboot to restore GPU to detectable state

Test 3: GPU Persistence Mode - FAIL

Attempted mitigation by enabling GPU persistence mode (`nvidia-smi -pm 1`) via a privileged init container before starting compute workloads. This is a standard workaround for intermittent Xid 79 errors.

Result: Persistence mode did not prevent the failure. GPU still fell off the bus under load.

Test 4: gpu-burn Stress Test - FAIL

Ran gpu-burn (CUDA stress test) to apply sustained compute load directly to the GPU.

Result: GPU failed with **Xid 79** and **PCIe RxErr** (PCIe receive errors), confirming the failure is hardware-related and occurs under any compute load, not specific to the Ollama application.

Error Details

NVIDIA Xid 79 - "GPU has fallen off the bus"

This is a critical hardware error indicating the GPU has lost communication with the host system over the PCIe bus. Common causes include:

- Defective GPU card
- Failing PCIe slot or riser
- Power delivery issues
- Thermal throttling leading to shutdown

PCIe RxErr (Receive Errors)

PCIe receive errors were logged alongside the Xid 79 events, indicating data corruption or signal integrity problems on the PCIe bus. This further supports a hardware-level failure.

Troubleshooting Steps Taken

| Step | Action | Result |
|------|---|--|
| 1 | Verified NVIDIA driver installation (570.211.01) | Driver loads correctly |
| 2 | Verified NVIDIA container toolkit and device plugin | GPU detected, allocatable in Kubernetes |
| 3 | Tested nvidia-smi at idle | GPU responds normally |
| 4 | Ran LLM inference workload | GPU crashes with Xid 79 |
| 5 | Enabled GPU persistence mode | Did not resolve issue |
| 6 | Rebooted host node | GPU recovers temporarily, fails again under load |
| 7 | Ran gpu-burn stress test | Confirmed failure: Xid 79 + PCIe RxErr |
| 8 | Switched all workloads to CPU-only mode | Workaround successful, services restored |

Timeline

| Date | Event |
|--------------|--|
| Feb 12, 2026 | GPU deployed with Ollama LLM server, initial operation successful |
| Feb 15, 2026 | GPU monitoring dashboards and temperature alerts configured |
| Feb 19, 2026 | GPU failures began; Xid 79 errors observed under load |
| Feb 19, 2026 | GPU persistence mode attempted as fix - unsuccessful |
| Feb 19, 2026 | gpu-burn stress test confirmed GPU hardware failure |
| Feb 19, 2026 | All workloads switched to CPU-only mode |
| Feb 20, 2026 | nvidia-smi confirms GPU still responds at idle but cannot sustain load |

Conclusion

The NVIDIA Tesla P4 GPU is defective. It passes basic detection and idle queries but **fails consistently under any compute workload** with Xid 79 (GPU fallen off the bus) and PCIe receive errors. Multiple mitigation attempts (persistence mode, node reboots) were unsuccessful. The failure pattern is consistent with a hardware defect in the GPU card itself. The GPU should be replaced under warranty/return.